

# COSC343: Artificial Intelligence

## Lecture 8 : Non-linear regression and Optimisation

Lech Szymanski

Dept. of Computer Science, University of Otago

## In today's lecture

- Error surface
- Iterative parameter search
- Non-linear regression and multiple minima
- Optimisation techniques:
  - Random walk
  - Steepest gradient
  - Simulated annealing

# Recall: Least squares and linear regression

The least squares parameters that minimise  $J = \frac{1}{2} \sum_i e_i^2$ , where  $e_i = y_i - \tilde{y}_i$ , can be found by solving:

$$\frac{dJ}{d\mathbf{w}} = \sum_i \frac{de_i}{d\mathbf{w}} e_i = 0$$

Linear regression is consistent only when appropriate choice of base functions is made

For models that are linear in parameters,  $y = f(\mathbf{x}, \mathbf{w}) = \mathbf{w}^T \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_U(\mathbf{x}) \end{bmatrix}$ , where  $\mathbf{w}^T = [w_1 \ \dots \ w_U]$ , there is a closed form solution:

$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F}\tilde{\mathbf{y}}^T, \text{ where}$$

Matrix inversion is computationally intensive

This matrix is sometimes singular

$$\mathbf{F} = \begin{bmatrix} f_1(\mathbf{x}_1) & \dots & f_1(\mathbf{x}_N) \\ \vdots & \ddots & \vdots \\ f_U(\mathbf{x}_1) & \dots & f_U(\mathbf{x}_N) \end{bmatrix} \text{ and } \tilde{\mathbf{y}} = [\tilde{y}_1 \ \dots \ \tilde{y}_N]$$

# Recall: Least squares and regression

The least squares parameters that minimise  $J = \frac{1}{2} \sum_i e_i^2$ , where  $e_i = y_i - \tilde{y}_i$ , can be found by solving:

$$\frac{dJ}{d\mathbf{w}} = \sum_i \frac{de_i}{d\mathbf{w}} e_i = 0$$

**Can appropriate parameters be found without the closed form equation?**


$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F}\tilde{\mathbf{y}}^T$$

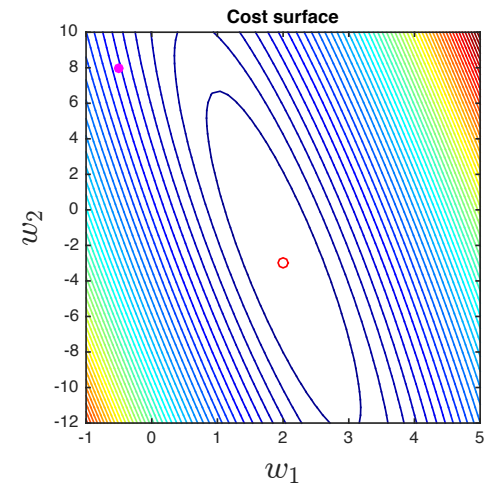

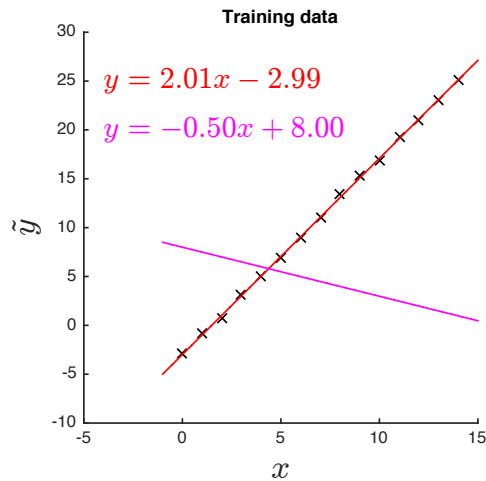
# Space of possible solutions

Let's think about the space of possible values that  $\mathbf{w}$  can take, and the corresponding cost

$$J = \frac{1}{2} \sum_i e_i^2 \text{ for some hypothesis } f(\mathbf{x}, \mathbf{w}) .$$

- The space of possible solutions is a U-dimensional **state-space** (where U is the number of parameters in the model);
- A **cost function** maps each point in the state space to a real number;
- The evaluations of all points in the state space can be visualised as a **state-space landscape** (in U+1 dimensions)

# An example: fitting a line



$$y_i = w_1 x_i + w_2$$

$$J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F} \tilde{\mathbf{y}}^T$$
$$w_1 = 2.01, w_2 = -2.99$$

closed form solution

Epoch0 :  $w_1 = -0.50, w_2 = 8.00$

initial guess

# Parameter search

1. Make an initial guess of the parameter values  $\mathbf{w}_0$  (often *random*) and compute the cost

$$J_0 = \frac{1}{2} \sum_i (f(\mathbf{x}_i, \mathbf{w}_0) - \tilde{y}_i)^2$$

2. Update the parameters and compute the new cost

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \Delta \mathbf{w}_t$$

Learning rate parameter  $\alpha$       Change in parameter values  $\Delta \mathbf{w}_t$

$$J_t = \frac{1}{2} \sum_i (f(\mathbf{x}_i, \mathbf{w}_t) - \tilde{y}_i)^2$$

3. Go back to step 2

# Random walk

Pick the weight update at random

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \dots \Delta w_{Ut}]^T, \text{ where}$$

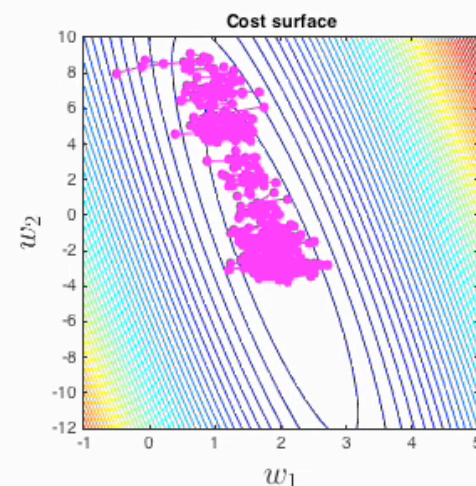


$$\Delta w_{jt} \sim \mathcal{N}(0, 1)$$

Parameter update is a random variable  
(in this case with a normal distribution)

keep the update  $\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \Delta \mathbf{w}_t$  if  
 $J_{t+1} < J_t$



# An example: fitting a line with random walk



$$y_i = w_1 x_i + w_2$$

$$J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F}\tilde{\mathbf{y}}^T$$

$$w_1 = 2.01, w_2 = -2.99$$

$$J = 0.18$$

Random walk:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$

$$\Delta w_{jt} \sim \mathcal{N}(0, 1)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + 0.50 \Delta \mathbf{w}_t$$

$$\text{Epoch 500 : } w_1 = 2.28, w_2 = -2.90$$




$$J = 41.94$$

# Learning rate parameter

Learning parameter controls how big the *update steps* are when parameters are updated.

- Small  $\alpha$  results in smaller steps: more ordered and direct path towards the minimum, but it takes longer to get there
- Large  $\alpha$  results in bigger steps: faster convergence, but less direct path and more chance of overshooting the minimum

# An example: random walk with smaller $\alpha$



$$y_i = w_1 x_i + w_2$$

$$J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F}\tilde{\mathbf{y}}^T$$

$$w_1 = 2.01, w_2 = -2.99$$

$$J = 0.18$$

Random walk:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$




$$\Delta w_{jt} \sim \mathcal{N}(0, 1)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + 0.200 \Delta \mathbf{w}_t$$

$$\text{Epoch 500 : } w_1 = 1.43, w_2 = 3.08$$

$$J = 77.51$$

# An example: random walk with larger $\alpha$



$$y_i = w_1 x_i + w_2$$

$$J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F}\tilde{\mathbf{y}}^T$$

$$w_1 = 2.01, w_2 = -2.99$$

$$J = 0.18$$

Random walk:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$

$$\Delta w_{jt} \sim \mathcal{N}(0, 1)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + 0.800 \Delta \mathbf{w}_t$$

$$\text{Epoch 500 : } w_1 = 2.68, w_2 = -2.35 \quad J = 280.02$$

# Steepest gradient descent

The weight update is the negative gradient of the cost function with respect to the parameters

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \Delta \mathbf{w}_t$$

Gradient is the positive slope (going up) of the cost surface at  $\mathbf{w}_t$

and

$$\Delta \mathbf{w}_t = -\frac{dJ}{d\mathbf{w}} = - \left[ \frac{\partial J}{\partial w_{1t}} \quad \dots \quad \frac{\partial J}{\partial w_{Ut}} \right]^T$$

- Also referred to as “steepest gradient ascent” or “hill climbing” if the objective function is being maximised

# An example: fitting a line with steepest gradient descent

$$\Delta \mathbf{w}_t = -\frac{dJ}{d\mathbf{w}} = - \left[ \frac{\partial J}{\partial w_{1t}} \quad \cdots \quad \frac{\partial J}{\partial w_{Ut}} \right]^T$$

Since  $J = \sum_i e_i^2$ ,

Given:  $y_i = w_1 x_i + w_2$ ,

then  $\frac{\partial J}{\partial w_j} = \sum_i \frac{\partial e_i}{\partial w_j} e_i$ .

Output derivatives are:

Cost derivatives are:

Since  $e_i = y_i - \tilde{y}_i$ ,

then  $\frac{\partial e_i}{\partial w_j} = \frac{\partial y_i}{\partial w_j}$ .

$$\frac{\partial y_i}{\partial w_1} = x_i$$

$$\frac{\partial J}{\partial w_1} = \sum_i x_i e_i$$

$$\frac{\partial y_i}{\partial w_2} = 1$$




$$\frac{\partial J}{\partial w_2} = \sum_i e_i$$

Since  $y_i = \sum_j w_j f_j(\mathbf{x}_i)$ ,

then  $\frac{\partial y_i}{\partial w_j} = f_j(\mathbf{x}_i)$ .

$j$  - index over number of weights  
 $i$  - index over number of training samples

# An example: fitting a line with steepest gradient descent



$$y_i = w_1 x_i + w_2$$

$$J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

$$\mathbf{w} = (\mathbf{F}\mathbf{F}^T)^{-1} \mathbf{F}\tilde{\mathbf{y}}^T$$

$$w_1 = 2.01, w_2 = -2.99$$

$$J = 0.18$$

Steepest gradient descent:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$

$$\Delta w_{1t} = -\sum_i x_i (y_i - \tilde{y}_i), \quad \Delta w_{2t} = -\sum_i (y_i - \tilde{y}_i)$$

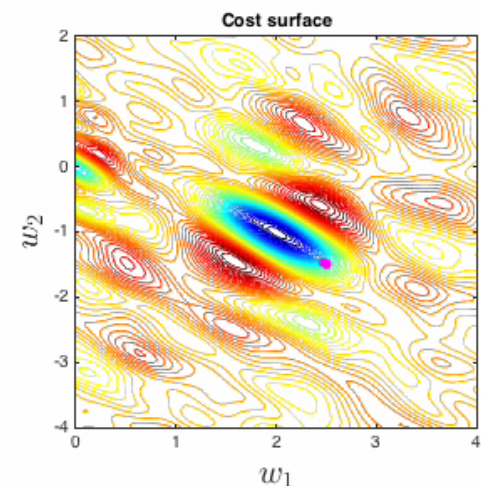


$$\mathbf{w}_{t+1} = \mathbf{w}_t + \frac{0.01}{N} \Delta \mathbf{w}_t$$

$$\text{Epoch 1 : } w_1 = 0.43, w_2 = 8.06$$

$$J = 349.93$$



# An example: non-linear regression



$$y_i = \sin(w_1x_{1i} + w_2x_{2i}) \quad J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

Steepest gradient descent:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$

$$\Delta w_{1t} = -\sum_i x_{1i} \cos(y_i)(y_i - \tilde{y}_i), \quad \Delta w_{2t} = -\sum_i x_{2i} \cos(y_i)(y_i - \tilde{y}_i)$$




$$\mathbf{w}_{t+1} = \mathbf{w}_t + \frac{0.10}{N} \Delta \mathbf{w}_t$$

Epoch1 :  $w_1 = 2.46, w_2 = -1.45$

$J = 78.00$



# An example: non-linear regression



$$y_i = \sin(w_1x_{1i} + w_2x_{2i}) \quad J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

Steepest gradient descent:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$


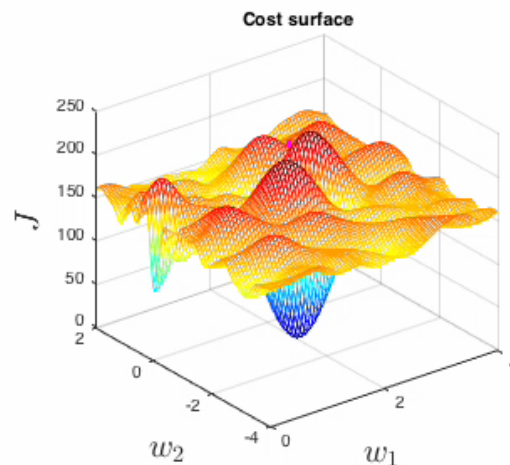

$$\Delta w_{1t} = -\sum_i x_{1i} \cos(y_i)(y_i - \tilde{y}_i), \quad \Delta w_{2t} = -\sum_i x_{2i} \cos(y_i)(y_i - \tilde{y}_i)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \frac{0.10}{N} \Delta \mathbf{w}_t$$

Epoch1 :  $w_1 = 1.03, w_2 = -1.79$

$J = 166.78$

# An example: non-linear regression



$$y_i = \sin(w_1x_{1i} + w_2x_{2i}) \quad J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

Steepest gradient descent:

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \ \Delta w_{2t}]^T$$

$$\Delta w_{1t} = -\sum_i x_{1i} \cos(y_i)(y_i - \tilde{y}_i), \quad \Delta w_{2t} = -\sum_i x_{2i} \cos(y_i)(y_i - \tilde{y}_i)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \frac{0.10}{N} \Delta \mathbf{w}_t$$

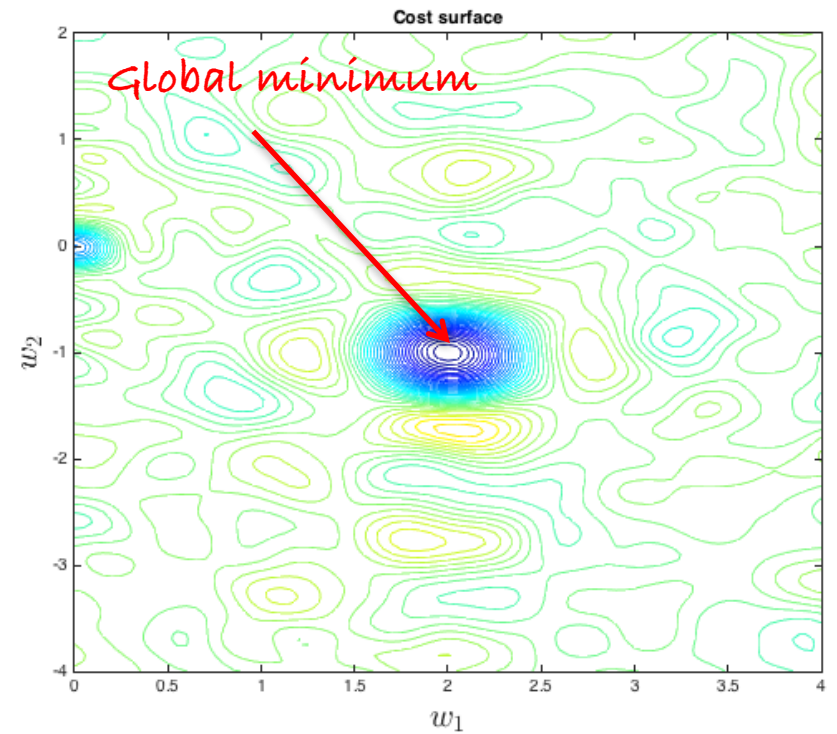

Epoch1 :  $w_1 = 2.97, w_2 = 1.10$

$J = 165.17$

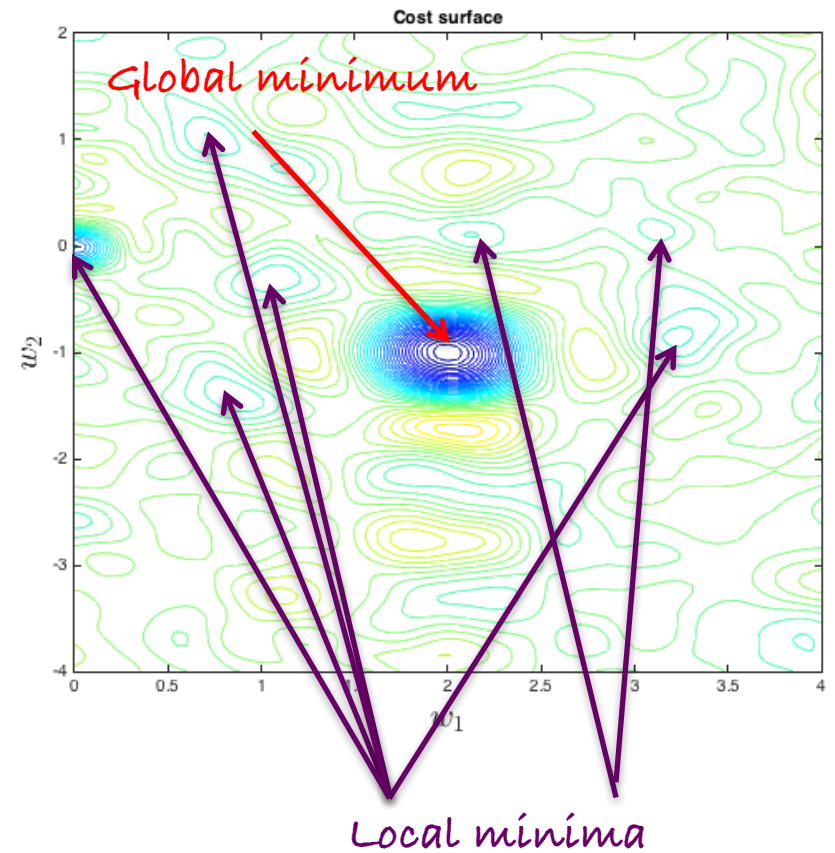

# Local minima

- Cost functions in non-linear optimisation are not necessarily convex
  - The cost surface may not correspond to a one giant valley, but a series of valleys separated by high cost ridges
- **Global minimum** – the state of the system that gives lowest possible cost
  - Best solution for the chosen model
- **Local minimum** – the minimum closest to the current state
  - May not be the best solution
  - Steepest gradient always points towards the local minimum, and as a result, the outcome of the training is highly dependent on the starting value of the parameters

# Local minima



# Local minima



# Simulated annealing


Pick the weight update at random

$$\Delta \mathbf{w}_t = [\Delta w_{1t} \dots \Delta w_{Ut}]^T, \text{ where } \Delta w_{jt} \sim \mathcal{N}(0, 1)$$

keep the update  $\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \Delta \mathbf{w}_t$  with probability

$$P(\mathbf{w}_{t+1} | \Delta \mathbf{w}_t) = \frac{1}{1 + e^{-\frac{\Delta J}{T_t}}}, \text{ where}$$

Reduction in cost


$$\Delta J = J_t - J_{t+1}$$

$$T_t = T_0 e^{-tT_c}$$




Temperature



# Simulated annealing: probability of accepting the update

$$P(\mathbf{w}_{t+1}|\Delta\mathbf{w}_t) = \frac{1}{1 + e^{-\frac{\Delta J}{T_t}}}$$

- High temperature increases the chance of accepting an update that results in higher cost
- Low temperature reduces the change of accepting an update that results in higher cost
- Start training with high temperature
  - More energy to jump out of the current minimum
- End training with low temperature
  - No energy to jump out of the current minimum









# Simulated annealing: cooling schedule

Starting temperature




$$T_t = T_0 e^{-tT_c}$$

Rate of cooling





# An example: non-linear regression with simulated annealing



$$y_i = \sin(w_1 x_{1i} + w_2 x_{2i})$$

$$J = \frac{1}{2} \sum_i (y_i - \tilde{y}_i)^2$$

Simulated annealing:



$$\Delta w_{jt} \sim \mathcal{N}(0, 1)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \frac{0.20}{N} \Delta \mathbf{w}_t$$

$$P(\mathbf{w}_{t+1} | \Delta \mathbf{w}_t) = \frac{1}{1 + e^{-(J_t - J_{t+1})/T_t}}, T_t = 30e^{-t/1000}$$

Epoch1 :  $w_1 = 1.00, w_2 = -1.80$

$J = 162.50$



# Summary and reading

- Learning as parameter search over the cost surface
- Models that are non-linear in parameters tend to have multiple minima
- Random walk – slow, can't guarantee where it goes
- Gradient descent – very fast, finds local minima
- Simulated annealing – theoretical promise to find global minimum (but slow and hard to get it to work in practice)

Reading for the lecture: AIMA Chapter 18 Section 4

Reading for next lecture: AIMA Chapter 18 Sections 7.1,7.2