

COSC343: Artificial Intelligence

Lecture 17: Problem solving and search

Alistair Knott

Dept. of Computer Science, University of Otago


In today's lecture

- The concept of state-space search
- Basic search algorithms

The concept of state-space search

Recap: in an **optimisation** task,


- There's a set of possible states;
- You can establish any state at any time;
- You don't know which state is best.



The concept of state-space search


In a **state-space search** task,

- There's a set of possible states;
- To get into a given state, you may have to go through *other* states.
- You *know* which state is best: but you don't know how to get there.



Actions and state transitions


Lots of problems can be formalised as state-space search tasks.



- The agent is in a given **start state**.
- The agent wants to reach some **goal state**.
- In each state, there is a set of **actions** the agent can perform.
 - In different states, there are different possible actions.
 - Each action gets you into a **new state**.

Look-ahead search

If agents can *represent* their world, and *reason* about it a little, they can solve state-space search tasks 'in their heads', without acting at all.




This is a distinctively human skill. It requires the ability

- to represent states of the world *other than the actual state*;
- to represent the results of actions *before they're executed*.


Recap from Lecture 1: types of agent

Goal-based agents can *search* for a way of achieving a desired situation.



Representation and reasoning

Our ability to reason about the consequences of our actions, in service of our goals, depends in large part on the **prefrontal cortex**.



Not fully developed in humans until age 21!

Formalising look-ahead search

Formally, a look-ahead search problem has three components:

- The agent's INITIAL STATE.
- A function $SUCCESSOR-FN(s)$ that takes a state s and returns the set of actions that are possible in s , along with the *new* state resulting from each action. (A set of $\langle action, result \rangle$ pairs)
- A GOAL TEST. (Either a set of states, or a Boolean fn on states.)

The first two of these define the **state space** of the problem.

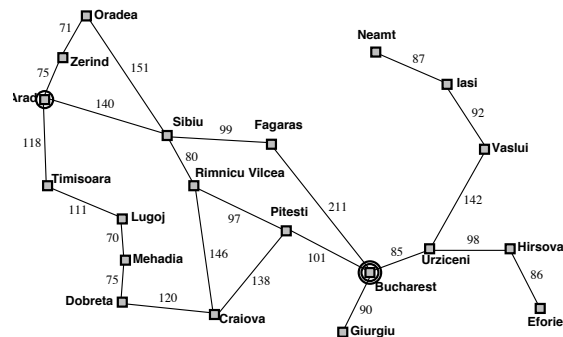
Example: Romania

Let's say the agent is on holiday in Romania, is currently in Arad, and needs to get to Bucharest.

- Let's model Romania as having 20 towns, linked by various roads.
- We can represent the ACTIONS and RESULT functions as a **state space graph**, in which the nodes are states (towns) and the arcs are actions (travelling down roads which connect towns together).
- The initial state is Arad; the goal state is Bucharest.


We can now define a formal search problem: is there a sequence of actions that takes the agent from Arad to Bucharest?

A state space graph for Romania



(Note: action costs are shown as labels on the arcs of the graph.)

Another example: The 8-puzzle



- states??
- actions??
- goal test??
- path cost??


Tree search algorithms

To solve a look-ahead search problem, we build a **search tree** which systematically explores the state space graph.

- We begin by creating a **root node** at the start state.
- Then we **expand** this node, by finding all the states we can get to from this state, and adding these nodes as children of the root node.
- This creates a **fringe** of new nodes. We now expand all of these new nodes in turn.
- Each time we expand a node, we add the new nodes to the fringe.
- Before we expand a node, we check to see if it's the goal state.

A search tree

Here's a simple search tree, whose root node is A.



Nodes and states

Don't confuse *nodes* in the search tree with *states* in the state-space graph!

- A **state** is a (representation of a) possible state of the world.
- A **node** is the computer's record of a state as it is encountered during a systematic search of the state space.

A node is a data structure that has several attributes.

- It has a **parent** node and a set of **successor** nodes.
- Since each node represents a state, one of its attributes is a **state**.
- We also store the **action** which got us to this state in the search.
- We can also store the **depth** of the node in the search tree.
- And we can add a **path cost**, storing the combined **step costs** of all the actions taken from the start state.

Implementation: general tree search

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
loop do
  if fringe is empty then return failure
  node ← REMOVE-FRONT(fringe)
  if GOAL-TEST(problem, STATE(node)) then return node
  fringe ← INSERTALL(EXPAND(node, problem), fringe)

function EXPAND(node, problem) returns a set of nodes
successors ← the empty set
for each action, result in SUCCESSOR-FN(problem, STATE[node]) do
  s ← a new NODE
  PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
  PATH-COST[s] ← PATH-COST[node] + STEP-COST(STATE[node], action,
result)
  DEPTH[s] ← DEPTH[node] + 1
  add s to successors
return successors
```

Search strategies

There are different ways of building the search tree—i.e. different **search strategies**.

- The strategy is determined by the order in which nodes in the fringe are expanded.
- If the search algorithm always picks the first node in the fringe to expand, then the strategy is determined by how new nodes are *added* to the fringe.

Evaluating a search strategy

Strategies are evaluated along the following dimensions:

- **completeness**—does it always find a solution if one exists?
- **time complexity**—number of nodes generated/expanded
- **space complexity**—maximum number of nodes in memory
- **optimality**—does it always find a least-cost solution?

Time and space complexity are measured in terms of:


- b —maximum branching factor of the search tree
- d —depth of the least-cost solution
- m —maximum depth of the state space (may be ∞)

Breadth-first search

In **breadth-first search**, we expand nodes 'row-by-row' in the search tree.

- Algorithm: put new successor nodes at the *end* of the fringe.

Q: What order would nodes be expanded in the following search tree?



Properties of breadth-first search

- **Complete??** Yes (if b is finite)
- **Time??** $1 + b + b^2 + b^3 + \dots + b^d + b(b^{d+1}) = O(b^{d+1})$
i.e., exponential in d
- **Space??** $O(b^{d+1})$ (keeps every node in the lowest ply of the tree in memory)
- **Optimal??** Yes (if cost = 1 per step); not optimal in general

Space is the big problem.

- Breadth-first search can easily generate nodes at 100MB/sec
- so 24hrs = 8640GB!

Uniform-cost search

Uniform-cost search is a variant of breadth-first search.

- We associate each node with a **path cost** g .
- When choosing a node to expand, always pick the node with lowest g .

Implementation: Keep the fringe ordered by path cost of nodes (with the lowest path cost first).

- **Complete??** Yes, if step cost $\geq \epsilon$
- **Time??** # of nodes with $g \leq$ cost of optimal solution, $O(b^{\lceil C^*/\epsilon \rceil})$ where C^* is the cost of the optimal solution
- **Space??** # of nodes with $g \leq$ cost of optimal solution, $O(b^{\lceil C^*/\epsilon \rceil})$
- **Optimal??** Yes—nodes expanded in increasing order of $g(n)$

Properties of depth-first search


- **Complete??** No: fails in infinite-depth spaces, spaces with loops.
 - You can modify the algorithm to avoid repeated states.
 - Then it's complete in finite spaces.
- **Time??** $O(b^m)$: terrible if m is much larger than d .
But if solutions are dense, may be much faster than breadth-first.
- **Space??** $O(bm)$, i.e., linear space!
- **Optimal??** No

Depth-first search

In **depth-first search**, we expand the most recent successor node we generated.

- Algorithm: put new successor nodes at the *front* of the fringe.


Q: What order would nodes be expanded in the following search tree?



Depth-limited search

Depth-limited search is a depth-first search with a depth limit of l . (I.e. nodes at depth l have no successors.)

If $l = 0$, we just search the start node.



Iterative deepening search

In **iterative deepening search**, we iteratively perform a depth-limited search in the tree, setting l first to 0, then to level 1, and so on.

```

function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution
  inputs: problem, a problem

  for depth  $\leftarrow$  0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
  end
    
```

Summary of search algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes
Time	b^{d+1}	$b^{\lceil C^*/\epsilon \rceil}$	b^m	b^l	b^d
Space	b^{d+1}	$b^{\lceil C^*/\epsilon \rceil}$	bm	bl	bd
Optimal?	Yes*	Yes	No	No	Yes*

Properties of iterative deepening search

- **Complete??** Yes
- **Time??** $(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- **Space??** $O(bd)$
- **Optimal??** Yes, if step cost = 1.
(The algorithm can be modified to explore uniform-cost tree.)

Q: How many nodes must be stored to search a tree with $b = 10$ and $d = 5$, if the solution is at the far right leaf?

$$N(\text{IDS}) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$


$$N(\text{BFS}) = 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,100$$

IDS does better because other nodes at depth d are not expanded.

Repeated states

Failure to detect repeated states can turn a linear problem into an exponential one!

For instance: the search space on the left is linear, but if we don't ignore repeated states, we produce a binary branching search tree.



Graph search

In **graph search**:

- We keep a list called *closed*, holding all the states we have encountered so far.
- We only add a node to the fringe if it's not in *closed*.

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
  end
```

Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored.
- There are several different uninformed search strategies.
- Iterative deepening is pretty much the best of these.
- Graph search can be exponentially more efficient than tree search.

Reading

The reading for today's lecture is AIMA Sections 3.1–3.4.

For next lecture: AIMA Section 3.5–3.6.